

HW 10

Rakshitha Panduranga

Problem 2 a

In [66]:

```
1 import numpy as np
2 import pandas as pd
3 from sklearn import svm
4 import matplotlib.pyplot as plt
5 from sklearn.model_selection import KFold
```

In [385]:

```
1 feature_tr = np.genfromtxt('feature_train.csv', delimiter=',');
2 label_tr = np.genfromtxt('label_train.csv', delimiter=',');
3 feature_te = np.genfromtxt('feature_test.csv', delimiter=',');
4 label_te = np.genfromtxt('label_test.csv', delimiter=',');
5
```

In [386]:

```
1 feature_tr
```

Out[386]:

```
array([[1.375e+01, 1.730e+00, 2.410e+00, ..., 1.150e+00, 2.900e+00,
       1.320e+03],
       [1.320e+01, 1.780e+00, 2.140e+00, ..., 1.050e+00, 3.400e+00,
       1.050e+03],
       [1.307e+01, 1.500e+00, 2.100e+00, ..., 1.180e+00, 2.690e+00,
       1.020e+03],
       ...,
       [1.253e+01, 5.510e+00, 2.640e+00, ..., 8.200e-01, 1.690e+00,
       5.150e+02],
       [1.288e+01, 2.990e+00, 2.400e+00, ..., 7.400e-01, 1.420e+00,
       5.300e+02],
       [1.373e+01, 4.360e+00, 2.260e+00, ..., 7.800e-01, 1.750e+00,
       5.200e+02]])
```

In [387]:

```
1 label_tr
```

Out[387]:

```
array([1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
       1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
       2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2.,
       2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2., 2.,
       3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3.,
       3., 3., 3., 3.])
```

In [388]:

```
1 label_tr = np.vstack(label_tr)
2 label_te = np.vstack(label_te)
```

In [389]:

```
1 train_data = feature_tr[:,0:2]
2 test_data = feature_te[:,0:2]
```

In [130]:

```
1 train_data_wlabels = np.column_stack((train_data,label_tr))
```

In [97]:

```
1 n_cv = 5
2 cv = StratifiedKFold(n_cv)
3 precomputed_folds = list(cv.split(train_data, label_tr))
```

In [98]:

1 precomputed_folds

Out[98]:

```
[((array([ 6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22,
       23, 24, 25, 26, 27, 28, 29, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46,
       47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63,
       64, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83, 84, 85,
       86, 87, 88]),
 array([ 0,  1,  2,  3,  4,  5, 30, 31, 32, 33, 34, 35, 36, 65, 66, 67, 68,
       69])), 
 (array([ 0,  1,  2,  3,  4,  5, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22,
       23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 44, 45, 46,
       47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63,
       64, 65, 66, 67, 68, 69, 75, 76, 77, 78, 79, 80, 81, 82, 83, 84, 85,
       86, 87, 88]),
 array([ 6,  7,  8,  9, 10, 11, 37, 38, 39, 40, 41, 42, 43, 70, 71, 72, 73,
       74])), 
 (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 18, 19, 20, 21, 22,
       23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39,
       40, 41, 42, 43, 51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63,
       64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 80, 81, 82, 83, 84, 85,
       86, 87, 88]),
 array([12, 13, 14, 15, 16, 17, 44, 45, 46, 47, 48, 49, 50, 75, 76, 77, 78,
       79])), 
 (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
       17, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39,
       40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 58, 59, 60, 61, 62, 63,
       64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 85,
       86, 87, 88]),
 array([18, 19, 20, 21, 22, 23, 51, 52, 53, 54, 55, 56, 57, 80, 81, 82, 83,
       84])), 
 (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
       17, 18, 19, 20, 21, 22, 23, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39,
       40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56,
       57, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80,
       81, 82, 83, 84])), 
 array([24, 25, 26, 27, 28, 29, 58, 59, 60, 61, 62, 63, 64, 85, 86, 87, 8
 8]))]
```

In [155]:

```

1 accuracy = [];
2 for train_index, test_index in cv.split(train_data, label_tr):
3     count = 0;
4     X_train, X_test = train_data[train_index], train_data[test_index];
5     clf = svm.SVC(C = 1,kernel='rbf',gamma = 1)
6     clf.fit(X_train,train_data_wlabels[train_index,2:3])
7     labels = clf.predict(X_test)
8     sample = train_data_wlabels[test_index,2:3]
9     for i in range(0,len(labels)):
10         if(labels[i] == sample[i]):
11             count = count+1;
12     accuracy.append((count/len(labels))*100);
13 average_acc = np.mean(accuracy);

```

C:\Users\Hp\Anaconda3\lib\site-packages\sklearn\utils\validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

In [156]:

```
1 accuracy
```

Out[156]:

```
[94.44444444444444,
 83.33333333333334,
 66.66666666666666,
 88.8888888888889,
 76.47058823529412]
```

In [308]:

```
1 print("The average cross-validation accuracy =",average_acc)
```

The average cross-validation accuracy = 81.96078431372548

Problem 2 (b)

In [237]:

```

1 Ny = 50;
2 Nc = 50;
3 y = np.logspace((-3),(3),Ny,base = 10)
4 c= np.logspace((-3),(3),Nc,base = 10)
```

In [292]:

1 c

Out[292]:

```
array([1.00000000e-03, 1.32571137e-03, 1.75751062e-03, 2.32995181e-03,
       3.08884360e-03, 4.09491506e-03, 5.42867544e-03, 7.19685673e-03,
       9.54095476e-03, 1.26485522e-02, 1.67683294e-02, 2.22299648e-02,
      2.94705170e-02, 3.90693994e-02, 5.17947468e-02, 6.86648845e-02,
      9.10298178e-02, 1.20679264e-01, 1.59985872e-01, 2.12095089e-01,
      2.81176870e-01, 3.72759372e-01, 4.94171336e-01, 6.55128557e-01,
     8.68511374e-01, 1.15139540e+00, 1.52641797e+00, 2.02358965e+00,
     2.68269580e+00, 3.55648031e+00, 4.71486636e+00, 6.25055193e+00,
     8.28642773e+00, 1.09854114e+01, 1.45634848e+01, 1.93069773e+01,
    2.55954792e+01, 3.39322177e+01, 4.49843267e+01, 5.96362332e+01,
    7.90604321e+01, 1.04811313e+02, 1.38949549e+02, 1.84206997e+02,
   2.44205309e+02, 3.23745754e+02, 4.29193426e+02, 5.68986603e+02,
   7.54312006e+02, 1.00000000e+03])
```

In [239]:

```
1 n_cv = 5;
2 def acc(C_value, g_value):
3     accuracy = [];
4     cv = StratifiedKFold(n_cv)
5     clf = svm.SVC(C = C_value,kernel='rbf',gamma = g_value)
6     for train_index, test_index in cv.split(train_data, label_tr):
7         count = 0;
8         X_train, X_test = train_data[train_index], train_data[test_index];
9         clf.fit(X_train,train_data_wlabels[train_index,2:3])
10        labels = clf.predict(X_test)
11        sample = train_data_wlabels[test_index,2:3]
12        for i in range(0,len(labels)):
13            if(labels[i] == sample[i]):
14                count = count+1;
15        accuracy.append((count/len(labels)));
16        average_acc = np.mean(accuracy);
17
18        sd_acc = np.std(accuracy);
19        return average_acc,sd_acc;
```

In [327]:

```

1 n_cv = 5;
2 def acc(C_value, g_value):
3     accuracy = [];
4     cv = StratifiedKFold(n_cv)
5     clf = svm.SVC(C = C_value,kernel='rbf',gamma = g_value)
6     for train_index, test_index in cv.split(train_data, label_tr):
7         count = 0;
8         X_train, X_test = train_data[train_index], train_data[test_index];
9         clf.fit(X_train,train_data_wlabels[train_index,2:3])
10        labels = clf.predict(X_test)
11        sample = train_data_wlabels[test_index,2:3]
12        for i in range(0,len(labels)):
13            if(labels[i] == sample[i]):
14                count = count+1;
15            accuracy.append((count/len(labels)));
16        average_acc = np.mean(accuracy);
17        squares = [];
18        M = len(accuracy);
19        for pm in range(0,M):
20            squares.append(((accuracy[pm])**2) - ((average_acc)**2))
21        sd_acc = (((1/(M-1))*np.sum(squares))**(1/2));
22
23 #sd_acc = np.std(accuracy);
24 return average_acc,sd_acc;

```

In [328]:

```

1 ACC = np.zeros((Ny,Nc));
2 SD = np.zeros((Ny,Nc));
3 for i in range(0,Ny):
4     for j in range(0,Nc):
5         #SD[i][j],ACC[i][j] = acc(values[i],values[j]);
6         a,b = acc(values[j],values[i])
7         SD[i][j] = b
8         ACC[i][j] = a
9
10

```

C:\Users\Hp\Anaconda3\lib\site-packages\sklearn\utils\validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

In [329]:

```
1 ACC
```

Out[329]:

```
array([[0.39346405, 0.39346405, 0.39346405, ..., 0.80915033, 0.80915033,
       0.82026144],
       [0.39346405, 0.39346405, 0.39346405, ..., 0.80915033, 0.80915033,
       0.80915033],
       [0.39346405, 0.39346405, 0.39346405, ..., 0.80915033, 0.80915033,
       0.82026144],
       ...,
       [0.39346405, 0.39346405, 0.39346405, ..., 0.43856209, 0.43856209,
       0.43856209],
       [0.39346405, 0.39346405, 0.39346405, ..., 0.42745098, 0.42745098,
       0.42745098],
       [0.39346405, 0.39346405, 0.39346405, ..., 0.42745098, 0.42745098,
       0.42745098]])
```

In [330]:

```
1 SD
```

Out[330]:

```
array([[0.01023038, 0.01023038, 0.01023038, ..., 0.08372011, 0.08372011,
       0.06056782],
       [0.01023038, 0.01023038, 0.01023038, ..., 0.08372011, 0.08372011,
       0.08372011],
       [0.01023038, 0.01023038, 0.01023038, ..., 0.08372011, 0.08372011,
       0.09920333],
       ...,
       [0.01023038, 0.01023038, 0.01023038, ..., 0.04942101, 0.04942101,
       0.04942101],
       [0.01023038, 0.01023038, 0.01023038, ..., 0.03678467, 0.03678467,
       0.03678467],
       [0.01023038, 0.01023038, 0.01023038, ..., 0.03678467, 0.03678467,
       0.03678467]])
```

In [331]:

```
1 best_fit = np.where(ACC == ACC.max())
```

In [332]:

```
1 best_fit = np.column_stack((best_fit[0],best_fit[1]))
```

In [333]:

```
1 best_fit
```

Out[333]:

```
array([[21, 37],
       [22, 32],
       [22, 33],
       [22, 34],
       [24, 28],
       [25, 27]], dtype=int64)
```

In [345]:

```
1 ACC.max()
```

Out[345]:

0.85359477124183

In [348]:

```
1 best_fit = np.where(ACC == ACC.max())
2 best_fit = np.column_stack((best_fit[0],best_fit[1]))
3 len(best_fit)
4
```

Out[348]:

6

In [355]:

```
1 SD_best = []
2 for i in range (0,len(best_fit)):
3     sp = best_fit[i,:]
4     SD_best.append(SD[sp[0]][sp[1]])
5 SDmin = np.min(SD_best)
6 for i in range(0,50):
7     for j in range(0,50):
8         if(ACC[i][j] == ACC.max() and SD[i][j] == SDmin):
9             iind = i;
10            jind = j;
11            break;
12 best_C = c[jind];
13 best_y = y[iind];
```

In [356]:

```
1 best_C
```

Out[356]:

8.286427728546842

In [357]:

```
1 best_y
```

Out[357]:

0.49417133613238334

In [358]:

```
1 SDmin
```

Out[358]:

0.08501130083463712

In [359]:

```
1 ACC.max()
```

Out[359]:

```
0.85359477124183
```

Problem 2 (c)

Without shuffling the data

In [365]:

```
1 ACC1 = np.zeros((Ny,Nc,20));
2 SD1 = np.zeros((Ny,Nc,20));
3 for t in range(0,20):
4     for i in range(0,Ny):
5         for j in range(0,Nc):
6             a,b = acc(values[j],values[i])
7             SD1[i][j][t] = b
8             ACC1[i][j][t] = a
9
10
```

```
C:\Users\Hp\Anaconda3\lib\site-packages\sklearn\utils\validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().  
y = column_or_1d(y, warn=True)
```

In [370]:

```
1 results = np.zeros((20,4))
```

In [371]:

```
1 for t in range(0,20):
2     ACC = ACC1[:, :, t]
3     SD = SD1[:, :, t]
4     best_fit = np.where(ACC == ACC.max())
5     best_fit = np.column_stack((best_fit[0], best_fit[1]))
6     SD_best = []
7     for i in range(0, len(best_fit)):
8         sp = best_fit[i, :]
9         SD_best.append(SD[sp[0]][sp[1]])
10    SDmin = np.min(SD_best)
11    for i in range(0, 50):
12        for j in range(0, 50):
13            if(ACC[i][j] == ACC.max() and SD[i][j] == SDmin):
14                iind = i;
15                jind = j;
16                break;
17    best_C = c[jind];
18    best_y = y[iind];
19    results[t, 0] = best_c;
20    results[t, 1] = best_y;
21    results[t, 2] = ACC.max();
22    results[t, 3] = SDmin;
23
```

In [372]:

1 results

Out[372]:

In [377]:

```

1 print("The final C = ",results[0][0]);
2 print("The final y = ",results[0][1]);
3 print("The final mean = ",results[0][2]);
4 print("The final SD = ",results[0][3]);
5

```

The final C = 33.9322177189533
 The final y = 0.28117686979742307
 The final mean = 0.8764705882352942
 The final SD = 0.024279968132030615

With Shuffling of data

In [375]:

```

1 n_cv = 5;
2 def acc1(C_value, g_value):
3     accuracy = [];
4     cv = StratifiedKFold(n_cv,shuffle = True)
5     clf = svm.SVC(C = C_value,kernel='rbf',gamma = g_value)
6     for train_index, test_index in cv.split(train_data, label_tr):
7         count = 0;
8         X_train, X_test = train_data[train_index], train_data[test_index];
9         clf.fit(X_train,train_data_wlabels[train_index,2:3])
10        labels = clf.predict(X_test)
11        sample = train_data_wlabels[test_index,2:3]
12        for i in range(0,len(labels)):
13            if(labels[i] == sample[i]):
14                count = count+1;
15        accuracy.append((count/len(labels)));
16        average_acc = np.mean(accuracy);
17        squares = [];
18        M = len(accuracy);
19        for pm in range(0,M):
20            squares.append(((accuracy[pm])**2) - ((average_acc)**2))
21        sd_acc = (((1/(M-1))*np.sum(squares))**((1/2)));
22
23        #sd_acc = np.std(accuracy);
24    return average_acc,sd_acc;

```

In [376]:

```

1 ACC1 = np.zeros((Ny,Nc,20));
2 SD1 = np.zeros((Ny,Nc,20));
3 for t in range(0,20):
4     for i in range(0,Ny):
5         for j in range(0,Nc):
6             a,b = acc1(values[j],values[i])
7             SD1[i][j][t] = b
8             ACC1[i][j][t] = a
9 for t in range(0,20):
10    ACC = ACC1[:, :, t]
11    SD = SD1[:, :, t]
12    best_fit = np.where(ACC == ACC.max())
13    best_fit = np.column_stack((best_fit[0],best_fit[1]))
14    SD_best = []
15    for i in range (0,len(best_fit)):
16        sp = best_fit[i,:]
17        SD_best.append(SD[sp[0]][sp[1]])
18    SDmin = np.min(SD_best)
19    for i in range(0,50):
20        for j in range(0,50):
21            if(ACC[i][j] == ACC.max() and SD[i][j] == SDmin):
22                iind = i;
23                jind = j;
24                break;
25    best_C = c[jind];
26    best_y = y[iind];
27    results[t,0] = best_c;
28    results[t,1] = best_y;
29    results[t,2] = ACC.max();
30    results[t,3] = SDmin;

```

C:\Users\Hp\Anaconda3\lib\site-packages\sklearn\utils\validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
`y = column_or_1d(y, warn=True)

In [379]:

```
1 results
```

Out[379]:

```
array([[3.39322177e+01, 2.81176870e-01, 8.76470588e-01, 2.42799681e-02],
       [3.39322177e+01, 2.81176870e-01, 8.66666667e-01, 1.00921678e-01],
       [3.39322177e+01, 3.72759372e-01, 8.64705882e-01, 5.14952218e-02],
       [3.39322177e+01, 9.10298178e-02, 8.66666667e-01, 8.42541716e-02],
       [3.39322177e+01, 9.10298178e-02, 8.76470588e-01, 8.22335476e-02],
       [3.39322177e+01, 6.55128557e-01, 8.54901961e-01, 6.22032019e-02],
       [3.39322177e+01, 2.81176870e-01, 8.76470588e-01, 8.22335476e-02],
       [3.39322177e+01, 4.94171336e-01, 8.65359477e-01, 4.90413895e-02],
       [3.39322177e+01, 3.72759372e-01, 8.66013072e-01, 8.35093617e-02],
       [3.39322177e+01, 2.81176870e-01, 8.64705882e-01, 3.32948644e-02],
       [3.39322177e+01, 2.12095089e-01, 8.66013072e-01, 6.23489674e-02],
       [3.39322177e+01, 3.72759372e-01, 8.66013072e-01, 9.22877206e-02],
       [3.39322177e+01, 3.90693994e-02, 8.66013072e-01, 9.22877206e-02],
       [3.39322177e+01, 6.55128557e-01, 8.65359477e-01, 4.90413895e-02],
       [3.39322177e+01, 3.72759372e-01, 8.66666667e-01, 8.42541716e-02],
       [3.39322177e+01, 2.81176870e-01, 8.76470588e-01, 6.06295028e-02],
       [3.39322177e+01, 6.55128557e-01, 8.65359477e-01, 2.93572479e-02],
       [3.39322177e+01, 3.72759372e-01, 8.77777778e-01, 7.24355823e-02],
       [3.39322177e+01, 6.86648845e-02, 8.66666667e-01, 8.42541716e-02],
       [3.39322177e+01, 8.68511374e-01, 8.66013072e-01, 8.35093617e-02]])
```

In [380]:

```
1 max_accuracy = np.max(results[:,2])
```

In [381]:

```
1 max_accuracy
```

Out[381]:

```
0.8777777777777779
```

In [382]:

```
1 for i in range(0,len(results)):
2     if(results[i][2] == max_accuracy):
3         best_c = results[i][0];
4         best_y = results[i][1];
5         best_mean = results[i][2];
6         best_sd = results[i][3];
7
```

In [383]:

```

1 print("The final C = ",best_c);
2 print("The final y = ",best_y);
3 print("The final mean = ",best_mean);
4 print("The final SD = ",best_sd);

```

The final C = 33.9322177189533
 The final y = 0.3727593720314938
 The final mean = 0.8777777777777779
 The final SD = 0.07243558228002855

Problem 2 d

In [396]:

```

1 accuracy = 0;
2 count = 0;
3 clf = svm.SVC(C = best_c,kernel='rbf',gamma = best_y)
4 clf.fit(train_data,label_tr)
5 labels = clf.predict(test_data)
6 for i in range(0,len(label_te)):
7     if(labels[i] == label_te[i]):
8         count = count+1;
9 accuracy= ((count/len(labels))*100);
10

```

C:\Users\Hp\Anaconda3\lib\site-packages\sklearn\utils\validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
 y = column_or_1d(y, warn=True)

In [397]:

```
1 accuracy
```

Out[397]:

78.65168539325843

Problem 1 a

In [398]:

```

1 from sklearn import svm
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt

```

In [399]:

```
1 dataset1_x = pd.DataFrame.as_matrix(pd.read_csv('train_x.csv'))
2 dataset1_y = pd.read_csv('train_y.csv')
3
4 dataset1_y = dataset1_y.values;
5 dataset1_y = dataset1_y.ravel();
6
7 dataset2_x = pd.DataFrame.as_matrix(pd.read_csv('train_x2.csv'))
8 dataset2_y = pd.read_csv('train_y2.csv')
9
10 dataset2_y = dataset2_y.values;
11 dataset2_y = dataset2_y.ravel();
```

C:\Users\Hp\Anaconda3\lib\site-packages\ipykernel_launcher.py:1: FutureWarning: Method .as_matrix will be removed in a future version. Use .values instead.

"""Entry point for launching an IPython kernel.

C:\Users\Hp\Anaconda3\lib\site-packages\ipykernel_launcher.py:7: FutureWarning: Method .as_matrix will be removed in a future version. Use .values instead.

```
import sys
```

In [404]:

```
1 def accuracy_classification(labels,length,dataset_y):
2     count = 0;
3     for i in range(0,length):
4         if(labels[i] == dataset_y[i]):
5             count += 1;
6     accuracy = (count/length)*100;
7     return accuracy;
```

In [405]:

```

1 def plotSVMBoundaries(training, label_train, classifier, support_vectors = []):
2
3     #Plot the decision boundaries and data points for minimum distance to
4     #class mean classifier
5     #
6     # training: traning data
7     # label_train: class lables correspond to training data
8     # classifier: sklearn classifier model, must have a predict() function
9     #
10    # Total number of classes
11    nclass = max(np.unique(label_train))
12
13    # Set the feature range for plotting
14    max_x = np.ceil(max(training[:, 0])) + 0.01
15    min_x = np.floor(min(training[:, 0])) - 0.01
16    max_y = np.ceil(max(training[:, 1])) + 0.01
17    min_y = np.floor(min(training[:, 1])) - 0.01
18
19    xrange = (min_x, max_x)
20    yrange = (min_y, max_y)
21
22    # step size for how finely you want to visualize the decision boundary.
23    inc = 0.005
24
25    # generate grid coordinates. this will be the basis of the decision
26    # boundary visualization.
27    (x, y) = np.meshgrid(np.arange(xrange[0], xrange[1]+inc/100, inc), np.arange(yrange[0], yrange[1]+inc/100, inc))
28
29    # size of the (x, y) image, which will also be the size of the
30    # decision boundary image that is used as the plot background.
31    image_size = x.shape
32    xy = np.hstack( (x.reshape(x.shape[0]*x.shape[1], 1, order='F'), y.reshape(y.shape[0]*y.shape[1], 1, order='F')))
33
34    # distance measure evaluations for each (x,y) pair.
35    pred_label = classifier.predict(xy)
36
37    # reshape the idx (which contains the class label) into an image.
38    decisionmap = pred_label.reshape(image_size, order='F')
39
40    #turn on interactive mode
41    plt.figure()
42    #plt.ion()
43
44    #show the image, give each coordinate a color according to its class label
45    plt.imshow(decisionmap, extent=[xrange[0], xrange[1], yrange[0], yrange[1]], origin='lower')
46
47    unique_labels = np.unique(label_train)
48    # plot the class training data.
49    plt.plot(training[label_train == unique_labels[0], 0], training[label_train == unique_labels[0], 1], 'o', color='red')
50    plt.plot(training[label_train == unique_labels[1], 0], training[label_train == unique_labels[1], 1], 'o', color='green')
51    if nclass == 3:
52        plt.plot(training[label_train == unique_labels[2], 0], training[label_train == unique_labels[2], 1], 'o', color='blue')
53
54    # include Legend for training data
55    if nclass == 3:
56        l = plt.legend(['Class 1', 'Class 2', 'Class 3'], loc=2)
57    else:
58        l = plt.legend(['Class 1', 'Class 2'], loc=2)
59    plt.gca().add_artist(l)

```

```
60
61     # plot support vectors
62     if len(support_vectors)>0:
63         sv_x = support_vectors[:, 0]
64         sv_y = support_vectors[:, 1]
65         plt.scatter(sv_x, sv_y, s = 100, c = 'blue')
66
67     plt.show()
```

In [412]:

```
1
2 def linear_SVM(C_value):
3     clf = svm.SVC(C = C_value,kernel='linear')
4     clf.fit(dataset1_x, dataset1_y)
5     labels = clf.predict(dataset1_x)
6     weight0 = clf.intercept_;
7     weight12 = clf.coef_;
8     weights = [];
9     weights.append(weight0);
10    weights.append(weight12);
11    sv = clf.support_vectors_;
12    decision_boundary = clf.decision_function;
13    return clf,sv,weights,labels,decision_boundary;
14
15 def RBF_Kernel(C_value):
16     clf = svm.SVC(C = C_value,kernel='rbf')
17     clf.fit(dataset2_x, dataset2_y)
18     labels = clf.predict(dataset2_x)
19     sv = clf.support_vectors_
20     return clf,sv,labels;
```

In [413]:

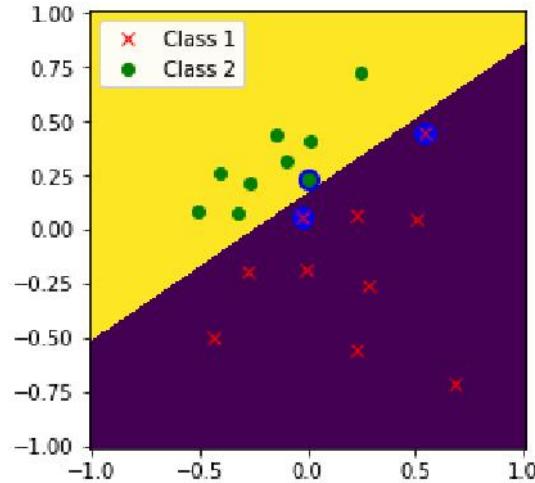
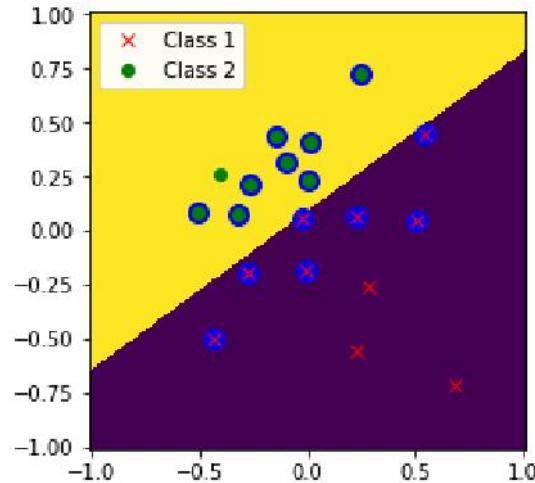
```

1 #1a
2 clf1,sv1,weights1,labels1,decision_boundary1 = np.array(linear_SVM(1));
3 accuracy1 = accuracy_classification(labels1,len(dataset1_y),dataset1_y);
4
5 clf100,sv100,weights100,labels100,decision_boundary100 = np.array(linear_SVM(100));
6 accuracy100 = accuracy_classification(labels100,len(dataset1_y),dataset1_y);
7
8 print("Accuracy when C = 1 is",accuracy1);
9 print("Accuracy when C = 100 is",accuracy100);
10
11 #plot
12 plot_1 = plotSVMBoundaries(dataset1_x,dataset1_y,clf1,sv1);
13 plot_100 = plotSVMBoundaries(dataset1_x,dataset1_y,clf100,sv100)
14
15

```

Accuracy when C = 1 is 100.0

Accuracy when C = 100 is 100.0



The value of C doesn't affect the classification in case of linear kernel. The accuracy for both C=1 and C=100 are the same.

problem 1 b

In [421]:

```
1 print("The weights when C = 100 are",weights100 );
2 print("The support vectors when C = 100 are",sv100);
3
4
```

```
The weights when C = 100 are [array([-1.79836766]), array([-7.11966384, 10.
40264821])]
The support vectors when C = 100 are [[-0.023855  0.06042
[ 0.54579   0.45029
[ 0.0064864  0.23394 ]]
```

Problem 1 c

The $g(x)$ is given by $= -1.798366 + (-7.11966384)x_1 + (10.40264821)x_2$ Yes, they lie on the boundary of the margin.

Problem 1 d

In [416]:

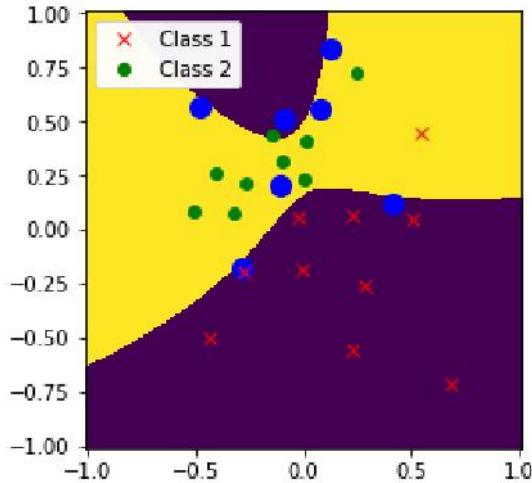
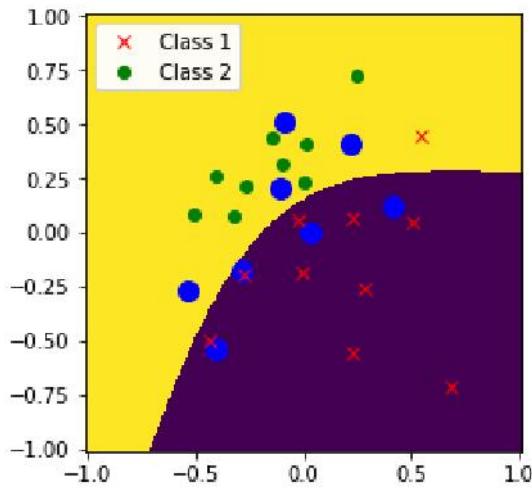
```

1 #1d
2 clf50,sv50,labels50 = np.array(RBF_Kernel(50));
3 accuracy50 = accuracy_classification(labels50,len(dataset2_y),dataset2_y);
4
5 clf5000,sv5000,labels5000 = np.array(RBF_Kernel(5000));
6 accuracy5000 = accuracy_classification(labels5000,len(dataset2_y),dataset2_y);
7
8 print("Accuracy when C = 50 is",accuracy50);
9 print("Accuracy when C = 5000 is",accuracy5000);
10
11 #plot
12 plot_50 = plotSVMBoundaries(dataset1_x,dataset1_y,clf50,sv50);
13 plot_5000 = plotSVMBoundaries(dataset1_x,dataset1_y,clf5000,sv5000)

```

Accuracy when C = 50 is 94.73684210526315

Accuracy when C = 5000 is 100.0



The decision boundaries are non- linear. the main differnece between the two decison boundaries are that tey are not in the same egration.

Problem 1 e

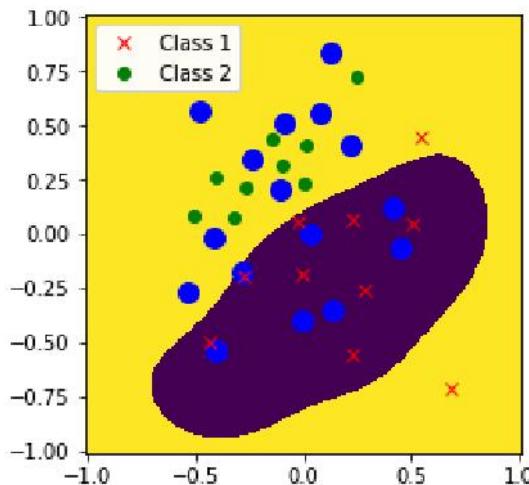
In [417]:

```
1 def RBF_Kernel1(gamma_value):  
2     clf = svm.SVC(kernel='rbf',gamma = gamma_value)  
3     clf.fit(dataset2_x, dataset2_y)  
4     labels = clf.predict(dataset2_x)  
5     sv = clf.support_vectors_  
6     return clf,sv,labels;  
7  
8  
9  
10  
11  
12  
13  
14  
15
```

In [418]:

```
1 clf10,sv10,labels10 = np.array(RBF_Kernel1(10));  
2 accuracy10 = accuracy_classification(labels10,len(dataset2_y),dataset2_y);  
3 print("Accuracy when gamma = 10 is",accuracy10);  
4 #plot  
5 plot_10 = plotSVMBoundaries(dataset1_x,dataset1_y,clf10,sv10);  
6
```

Accuracy when gamma = 10 is 94.73684210526315



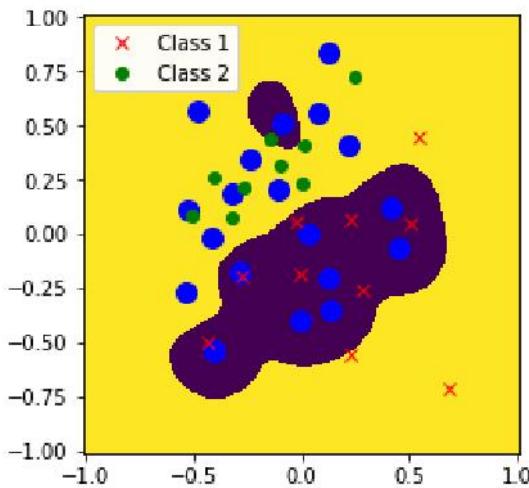
In [419]:

```

1 clf50,sv50,labels50 = np.array(RBF_Kernel1(50));
2 accuracy50 = accuracy_classification(labels50,len(dataset2_y),dataset2_y);
3 print("Accuracy when gamma = 50 is",accuracy50);
4 #plot
5 plot_50 = plotSVMBoundaries(dataset1_x,dataset1_y,clf50,sv50);
6

```

Accuracy when gamma = 50 is 100.0



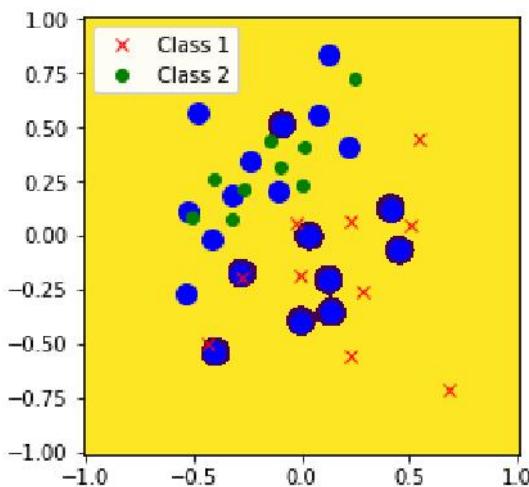
In [420]:

```

1 clf500,sv500,labels500 = np.array(RBF_Kernel1(500));
2 accuracy500 = accuracy_classification(labels500,len(dataset2_y),dataset2_y);
3 print("Accuracy when gamma = 500 is",accuracy500);
4 #plot
5 plot_500 = plotSVMBoundaries(dataset1_x,dataset1_y,clf500,sv500);
6

```

Accuracy when gamma = 500 is 100.0



With the increase in the gamma value, the decision regions are reducing drastically. Because of which, there is overfitting problem in the gamma= 500 plot.

In []:

1	
---	--